Write a Python program that implements gradient descent for minimizing
the least squares loss. As a stopping condition check for the objective (least squares error)
between the current and previous iteration. If the objective improves
by less than theta then you stop.

It should be run from the command line as follows:

python <your python program file> <data file name> <training file name>

Note that you need to convert label 0 to -1 from training label file since in least squares
model you have outputs of -1 or +1. But you should display your predictions as 0 instead of -1.